

Redis OM

❖ Introduction to Redis OM

- What is Redis OM?
 - Redis OM (pronounced "ohm") stands for **Object Mapping (and more)** for Redis. It's a family of high-level client libraries that provide powerful abstractions for working with Redis Stack, making it as easy as possible to use Redis as a document database with modern application frameworks.
 - Redis OM transforms Redis from a simple key-value store into a sophisticated document database by leveraging Redis Stack modules, particularly RedisJSON and RediSearch, to provide:
 - **Object Mapping:** Seamlessly map domain objects to Redis data structures
 - **Fluent Querying:** Query your data using language-native, type-safe APIs
 - **Automatic Indexing:** Efficient secondary indexes created automatically
 - **Full-Text Search:** Built-in search capabilities across your data
 - **Vector Similarity Search:** Modern AI/ML capabilities for semantic search
- Why Use Redis OM?
 - Traditional Redis usage requires developers to manually compose abstractions using core data structures. While powerful, this approach demands significant time and expertise. Redis OM solves this by providing:
 - **Developer Productivity:** Focus on business logic, not Redis internals
 - **Type Safety:** Compile-time guarantees and IDE support
 - **Performance:** Always-indexed queries for optimal performance
 - **Modern Features:** Vector search, geospatial queries, and full-text search
 - **Familiar APIs:** Language-idiomatic interfaces that feel natural

- Benefits and Use Cases

- Primary Benefits:

- Reduced boilerplate code for data access
 - Automatic query optimization through indexing
 - Type-safe operations with compile-time validation
 - Rich querying capabilities out of the box
 - High performance with Redis's in-memory architecture

- Common Use Cases:

- Document Databases: Store and query hierarchical data
 - Search Applications: Full-text search across content
 - Real-time Analytics: Fast aggregations and filtering
 - Recommendation Systems: Vector similarity for AI-powered recommendations
 - Geospatial Applications: Location-based queries and filtering
 - Caching with Querying: More than simple key-value caching

❖ Core Concepts

- Object Mapping and Schema Definition

- Redis OM automatically maps your domain objects to Redis data structures.

Objects can be stored as:

- Redis Hashes: For simple, flat structures
 - Redis JSON Documents: For complex, nested data (requires RedisJSON)

- Schema Definition Process:

- Annotate your domain classes with mapping annotations
 - Define field types and indexing strategies
 - Configure relationships and nested objects
 - Generate indexes automatically at runtime

- Indexing and Search Capabilities

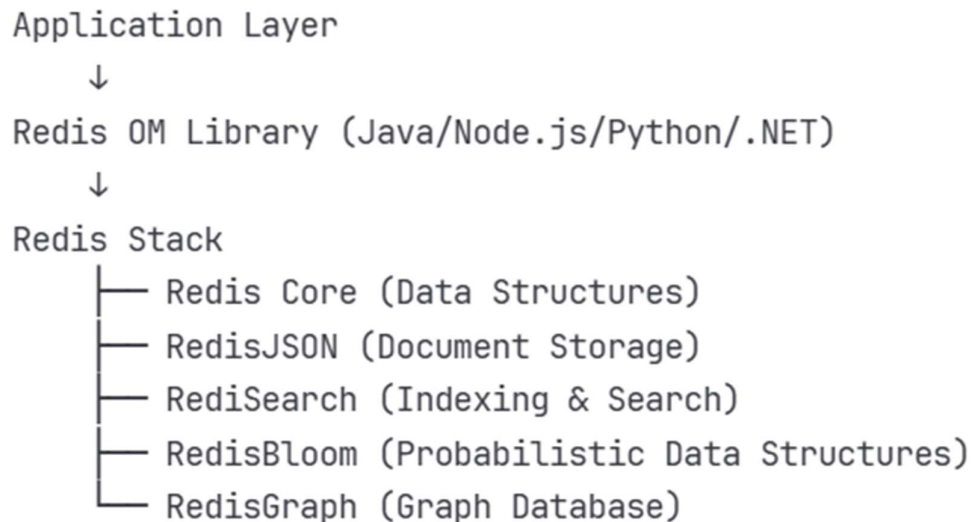
- Redis OM leverages RediSearch to provide powerful indexing:
- Index Types:
 - Tag Indexes: Exact-match searches (strings, enums)
 - Text Indexes: Full-text search with stemming and fuzzy matching
 - Numeric Indexes: Range queries and sorting
 - Geo Indexes: Geospatial queries within radius or bounding box
 - Vector Indexes: Similarity search for AI/ML applications
- Automatic Index Creation:
 - Indexes are created automatically based on annotations
 - Field-level control over indexing behavior
 - Support for nested object indexing
 - Composite indexes for multi-field queries

- CRUD Operations

- All Redis OM libraries provide consistent CRUD interfaces:
 - Create: Save new entities with auto-generated IDs
 - Read: Fetch by ID or query with complex criteria
 - Update: Modify existing entities with optimistic concurrency
 - Delete: Remove individual entities or bulk operations
- Advanced Operations:
 - Bulk operations for better performance
 - Upsert capabilities
 - Partial updates
 - Transactional operations where supported

❖ Architecture Overview

- How Redis OM Fits into Redis Ecosystem



- Client-Server Interactions

- Data Flow:

- **Entity Definition:** Domain objects defined with annotations
- **Index Creation:** Automatic schema generation and index creation
- **Data Persistence:** Objects serialized to JSON/Hash format
- **Query Processing:** High-level queries translated to Redis commands
- **Result Mapping:** Redis responses mapped back to domain objects

- Connection Management:

- Connection pooling for optimal performance
- Automatic retry and failover capabilities
- Support for Redis Cluster and Sentinel configurations
- SSL/TLS encryption support

❖ Supported Languages and SDKs

- Overview of Language-Specific Implementations

- Redis OM currently supports four major programming ecosystems:
 - Redis OM Spring (Java): Deep integration with Spring Framework
 - Redis OM Node.js: TypeScript-first with JavaScript support
 - Redis OM Python: Async/sync support with FastAPI integration
 - Redis OM .NET: LINQ support for C# developers

- Key Similarities Across SDKs

- Common Features:
 - Consistent annotation-based mapping
 - Repository pattern implementation
 - Fluent query APIs
 - Automatic index management
 - ULID-based ID generation
 - Vector similarity search support
- Shared Concepts:
 - Entity/Document annotations
 - Field-level indexing control
 - Relationship mapping
 - Query builders
 - Connection management

- Key Differences

- Language-Specific Adaptations:
 - Spring: Leverages Spring Data patterns and annotations
 - Node.js: TypeScript-first design with modern async/await
 - Python: Pydantic integration and FastAPI compatibility
 - .NET: LINQ query support and Entity Framework-like experience

❖ Common Features Across SDKs

- Data Modelling

- Entity Definition:

// Conceptual structure - syntax varies by language

@Document

class Product {

 @Id

 String id;

 @Indexed

 String name;

 @Searchable

 String description;

 @Indexed

 Double price;

 @Indexed

 Set<String> tags;

}

- Supported Field Types:

- Primitive types (string, number, boolean)
- Collections (arrays, sets, lists)
- Nested objects and documents
- Geospatial coordinates
- Date/time types
- Binary data (in some implementations)

- Querying

- Query Methods:

- Equality: Exact matches on indexed fields
 - Range: Numeric and date range queries
 - Full-Text: Search across text fields with ranking
 - Geospatial: Distance-based and bounding box queries
 - Tag Matching: Set membership and intersection queries
 - Vector Similarity: Semantic search with embeddings

- Query Building:

- Fluent API for complex query construction
 - Method chaining for multiple conditions
 - Sorting and pagination support
 - Aggregation capabilities

- Repository Pattern

- Standard Operations:

- save(entity): Persist entity to Redis
 - findById(id): Retrieve by primary key
 - findAll(): Get all entities of type
 - delete(entity): Remove entity
 - count(): Get total entity count

- Custom Queries:

- Method name-based query derivation
 - Annotation-based query definitions
 - Native Redis query support
 - Stream-based processing

- Indexing

- Automatic Index Management:

- Indexes created on application startup
 - Field-level index configuration
 - Composite index support
 - Index rebuilding capabilities

➤ Index Types Configuration:

- Text indexes with language-specific stemming
- Tag indexes for exact matching
- Numeric indexes with range support
- Geo indexes for location queries
- Vector indexes for similarity search